

Introduction to Data Science

An easy guide to simple data science and machine learning

Daniel Lee

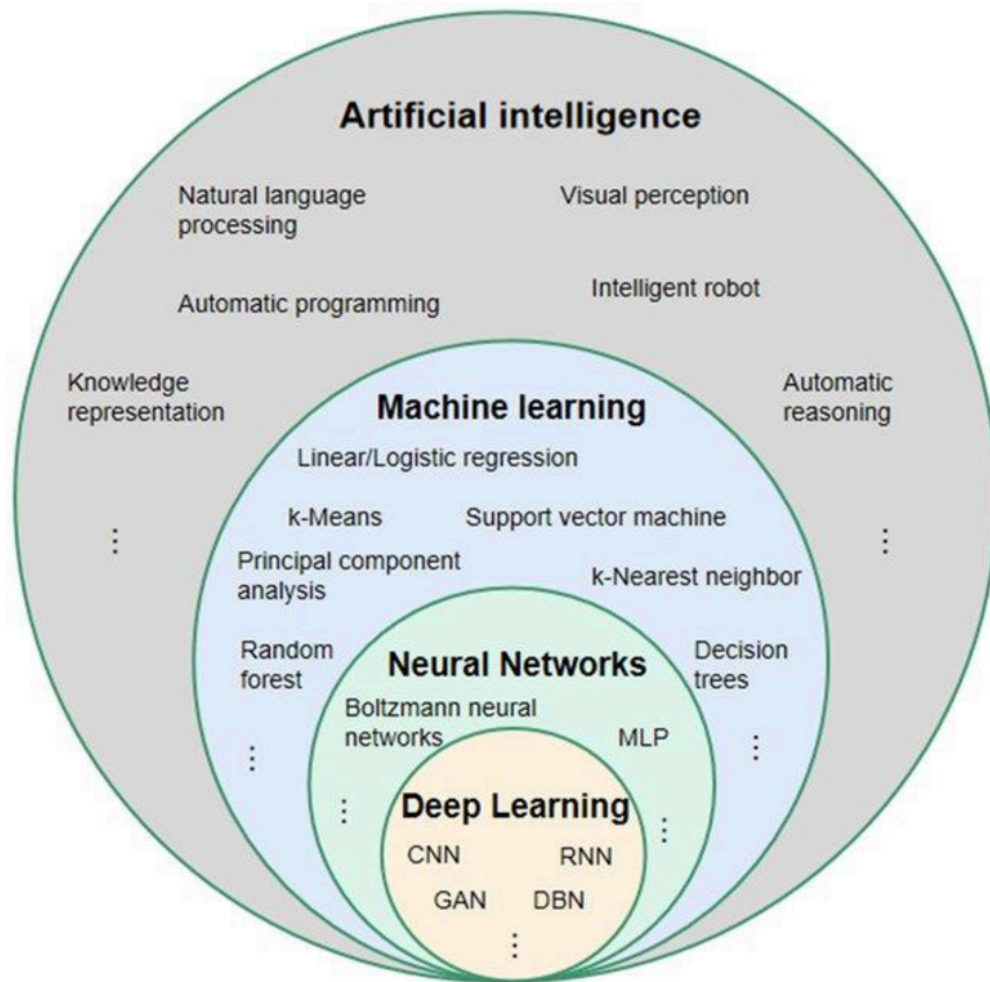
Contents

1. What is Data Science?
2. What is Machine Learning?
3. Principal Component Analysis (PCA)
4. Linear Regression
5. Logistic Regression

What is Data Science?



What is Machine Learning?




What is a Feature (X)?

A feature is an input variable used by a machine learning model. Features are also called "independent variables" or "predictors".

Example: In predicting Titanic survival,

- Features: class, sex, age, fare, etc.
- Target (Y): survived or not



Machine Learning as a Function




PCA (Principal Component Analysis)

PCA is a technique for reducing the dimensionality of data while preserving as much information as possible.


It finds new axes (principal components) that capture the largest **variance** in the data.



Principal Component Analysis (PCA) algorithm



Reduce data from 2D to 1D



Reduce data from 3D to 2D

PCA Example (Iris)

```
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

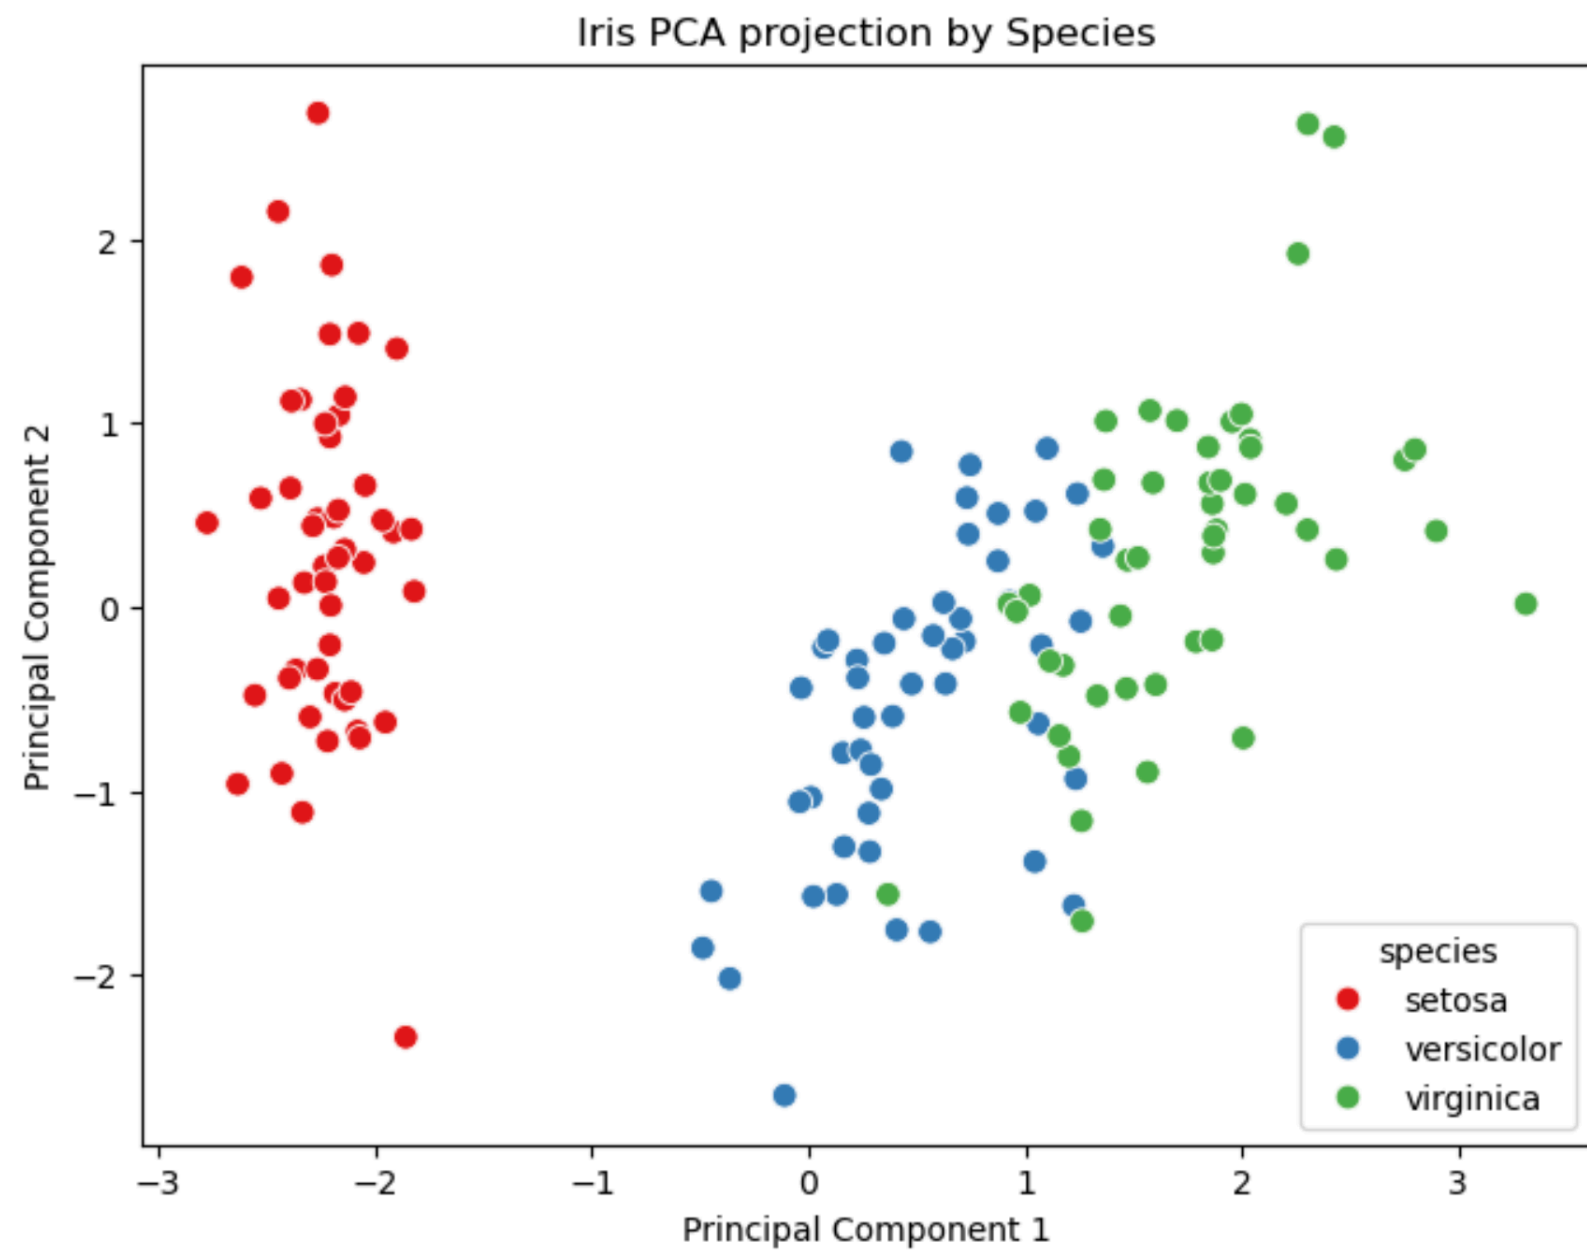
iris = sns.load_dataset("iris")
X = iris.drop("species", axis=1)
y = iris["species"]

X_scaled = StandardScaler().fit_transform(X)

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

iris['PC1'] = X_pca[:,0]
iris['PC2'] = X_pca[:,1]

plt.figure(figsize=(8,6))
sns.scatterplot(
    x='PC1', y='PC2',
    hue='species',
    data=iris,
    palette='Set1',
    s=60
)
plt.title('Iris PCA projection by Species')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

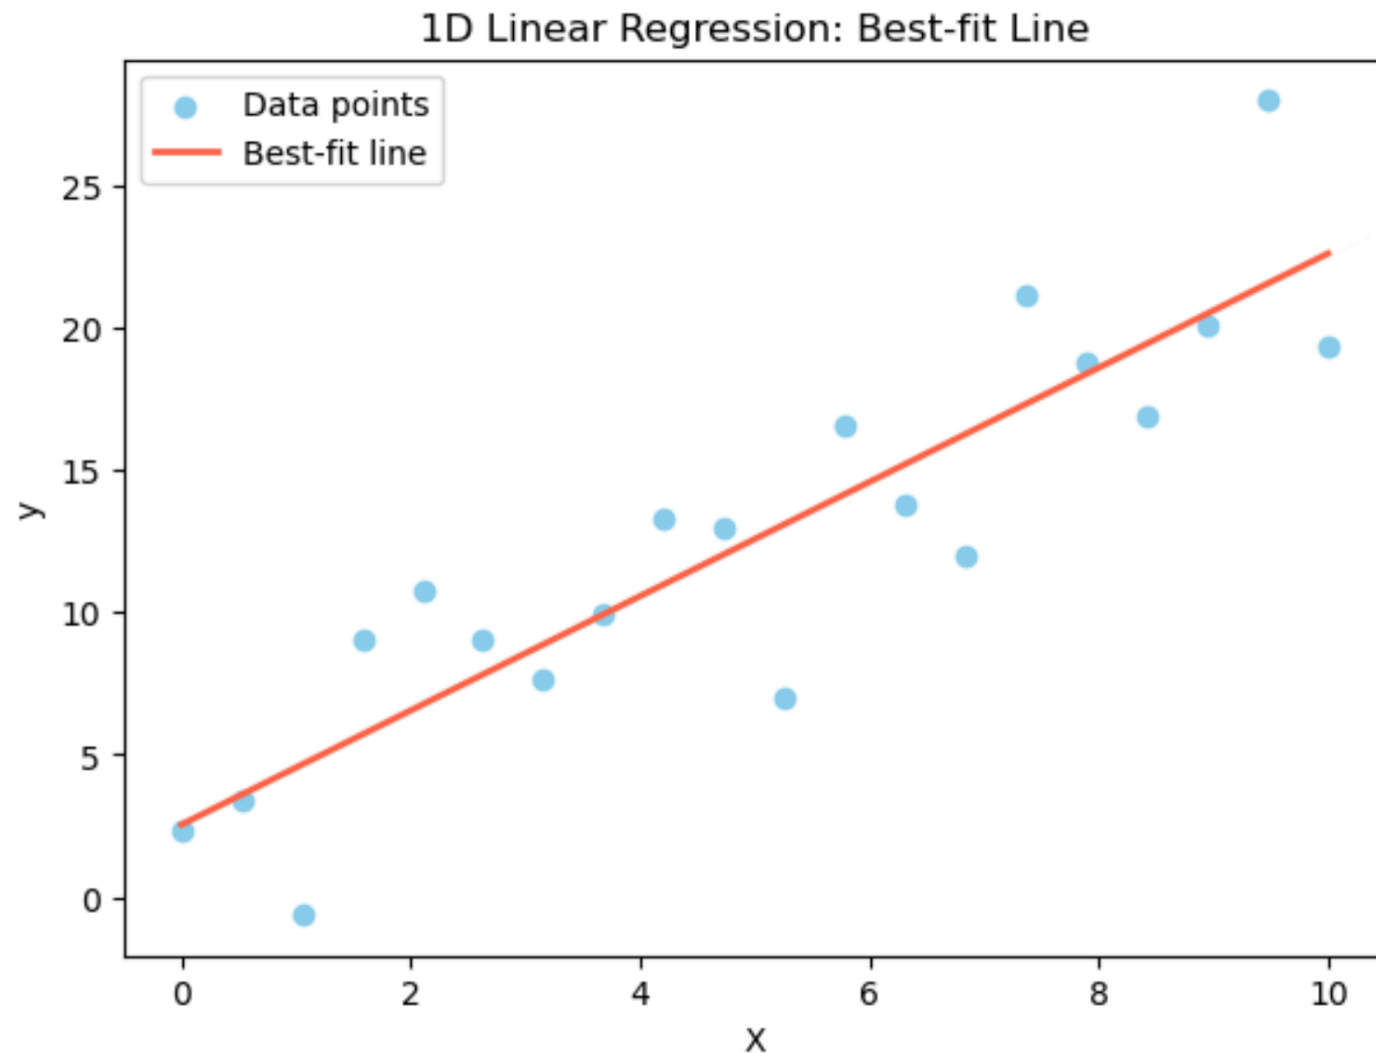


Linear Regression

Linear regression is a basic machine learning method that finds the best straight line to describe the relationship between two (or more) variables.


- It answers: "If X changes, how does Y change?"
- The line allows us to **predict** Y for any value of X.
- The best-fit line is as close as possible to all data points.

Linear Regression(1D)



Linear Regression(2D)

2D Linear Regression: Best-fit Plane



Get Correlations From Titanic

```
import seaborn as sns

df = sns.load_dataset("titanic")
df["sex"] = df["sex"].map({"male": 0, "female": 1})
corr = df.corr(numeric_only=True)
print(corr["survived"])
```

```
pclass      -0.338481
sex          0.543351
age         -0.077221
sibsp       -0.035322
parch        0.081629
fare         0.257307
adult_male  -0.557080
alone       -0.203367
Name: survived, dtype: float64
```

Titanic Example: Linear Regression

Linear regression models the relationship between features and a continuous target variable.

For Titanic, you might use regression to predict fare or age.


```

import seaborn as sns
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

df = sns.load_dataset("titanic")
df = df[["survived", "pclass", "sex", "age", "fare", "sibsp", "parch", "alone"]].dropna()
df["sex"] = df["sex"].map({"male": 0, "female": 1})

X = df[["pclass", "sex", "age", "fare", "alone"]]
y = df["survived"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred_reg = model.predict(X_test)

y_pred_class = (y_pred_reg >= 0.5).astype(int)

acc = accuracy_score(y_test, y_pred_class)
print("Linear Regression classification accuracy:", acc)

```

Titanic Example: Logistic Regression

Logistic regression is used for classification, such as predicting survival (yes/no) on the Titanic.

The model outputs the probability of the target being 1.

Logistic Regression Example

```
import seaborn as sns
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

df = sns.load_dataset("titanic")
df = df[["survived", "pclass", "sex", "age", "fare", "sibsp", "parch"]].dropna()
df["sex"] = df["sex"].map({"male": 0, "female": 1})
df["alive"] = df["alive"].map({"yes": 1, "no": 0})

X = df[["pclass", "sex", "age", "fare", "alone"]]
y = df["alive"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
```